

2022.1 Multicore Computing, Project #4

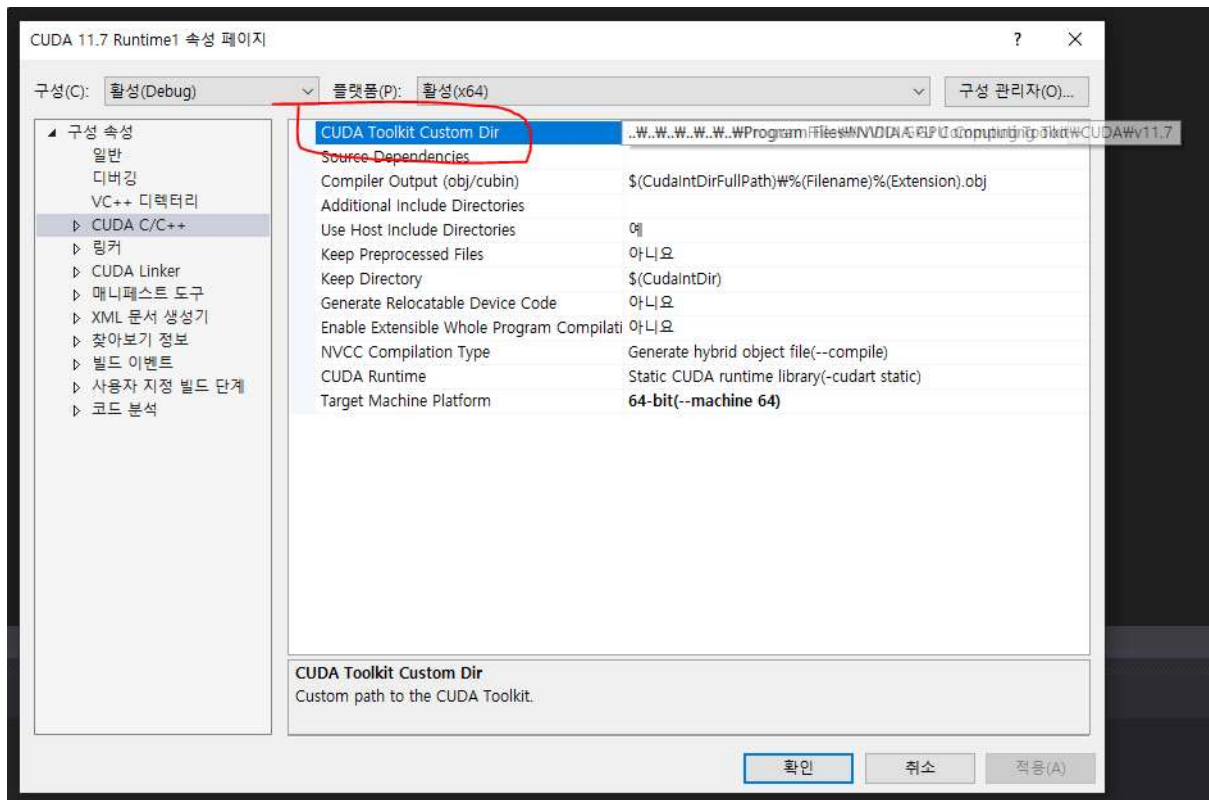
Problem 2

Document

소프트웨어학부

20176342 송민준

Set Windows SDK version to 10.0.17763.0



Set CUDA Toolkit Custom Dir like above picture

```
1>C:\Users\song\source\repos\CUDA 11.7 Runtime1\CUDA 11.7 Runtime1>..\..\..\..\Program
Files\NVIDIA GPU Computing Toolkit\CUDA\v11.7\bin\nvcc.exe -
gencode=arch=compute_52,code=sm_52,compute_52 --use-local-env -ccbin "C:\Program Files
(x86)\Microsoft Visual Studio\2017\Community\VC\Tools\MSVC\14.16.27023\bin\HostX86\x64" -x cu
-l"..\..\..\..\Program Files\NVIDIA GPU Computing Toolkit\CUDA\v11.7\include" -
l"..\..\..\..\Program Files\NVIDIA GPU Computing Toolkit\CUDA\v11.7\include" -G --
keep-dir x64\Debug -maxrregcount=0 --machine 64 --compile -cudart static -g -DWIN32 -DWIN64
-D_DEBUG -D_CONSOLE -D_MBCS -Xcompiler "/EHsc /W3 /nologo /Od /Fd" x64\Debug\vc141.pdb /FS /Zi
/RTC1 /MDd " -o "C:\Users\song\source\repos\CUDA 11.7 Runtime1\CUDA 11.7
Runtime1\x64\Debug\cuda_ray.cu.obj" "C:\Users\song\source\repos\CUDA 11.7 Runtime1\CUDA 11.7
Runtime1\cuda_ray.cu"
```

And compile project like above. Nvcc ~~~ thrust_ex.cu

(c) how to execute

'thrust_ex.exe' (no parameter)

Entire source code

```
#include "cuda_runtime.h"
#include "device_launch_parameters.h"

#include <stdio.h>

#include <stdio.h>
#include <string.h>
#include <stdlib.h>
#include <time.h>
#include <math.h>

#include <omp.h>
#include <cuda.h>

#include <thrust/host_vector.h>
#include <thrust/device_vector.h>
#include <thrust/transform.h>
#include <thrust/transform_reduce.h>

#define CUDA 0
#define OPENMP 1

#define rnd( x ) (x * rand() / RAND_MAX)
#define INF 2e10f

#define STEPS 1000000000
#define STEP 1/STEPS

struct saxpy_functor // define functor to calculate pie
{
    __host__ __device__
    double operator()(const int& x) const {
        double temp = (x + 0.5)*STEP;

        return (4.0 / (1.0 + temp * temp));
    }
};

int main()
{
    clock_t start, end;
```

```

start = clock();
double x, pi, sum = 0.0;

thrust::counting_iterator<int> a(0); //define counting iterator
double result = thrust::transform_reduce(a,a+STEPS,saxpy_functor(), 0.0,
thrust::plus<double>()); // reduce all result (add) with functor

pi = result * STEP; // calculate pi
printf("pi=%.8lf\n", pi);

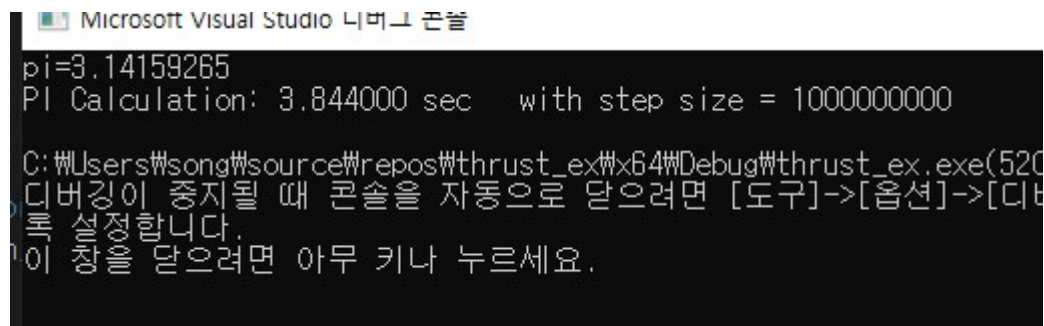
end = clock();//measure program execution time

printf("PI Calculation: %lf sec   with step size = %d \n", (double)(end -
start) / 1000.0,STEPS );
}

```

Program output result

- thrust_ex.cu



Microsoft Visual Studio 디버그 콘솔

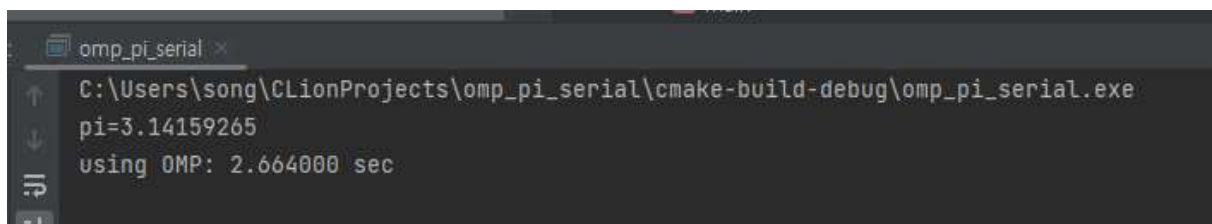
```

pi=3.14159265
PI Calculation: 3.844000 sec   with step size = 1000000000

C:\Users\song\source\repos\thrust_ex\x64\Debug\thrust_ex.exe(520:
디버깅이 중지될 때 콘솔을 자동으로 닫으려면 [도구]->[옵션]->[디버
로 설정합니다.
이 창을 닫으려면 아무 키나 누르세요.

```

- omp_pi_serial.c



```

omp_pi_serial x
C:\Users\song\CLionProjects\omp_pi_serial\cmake-build-debug\omp_pi_serial.exe
pi=3.14159265
using OMP: 2.664000 sec

```


Experimental results

Thrust_ex

Exec time(unit: sec)	
	3.844

Omp_pi_serial

Exec time(unit: sec)	1
	2.664



Interpretation/explanation

We calculated the pie value with n as 1 billion and calculated it through integration using openmp and 'cuda thrust'.

I think the reason why the performance of the open mp is better is because my cpu is the latest model, but I think it came out like this because the gpu is old.

And in the case of cuda thrust, the overhead of moving the value from the host vector to the device vector and moving the computation result back from the device vector to the host vector is too large, so if it is optimized, I think we can reduce the bus bottleneck and speed it up.